

Q1	Q2	Q3	Q4	Q5	Q6	Q7	Total (out of 25)
3.0	3.0	3.0	3.0	3.0	3.0	4.0	22.0

hl2248\_ipynb

May 3, 2023

```
[1]: import numpy as np

def syntheticY(n):
    """
    Synthetic rank-2 n x n matrix
    """
    U = np.vstack((np.arange(n)/n, (1-np.arange(n)/n)**2)).T
    W = np.vstack((np.arange(n)/n, (1-np.arange(n)/n)**2)).T
    return U @ W.T

def bernoulli(n,p):
    """
    Bernoulli sampling.
    Produces a set of index pairs I = {(i1,i2)} where
    i1,i2 = 0,...,n-1, and each pair is sampled with
    probability p
    """
    I1, I2 = np.meshgrid(np.arange(n), np.arange(n))
    I1 = I1.flatten()
    I2 = I2.flatten()
    I = []
    for i1,i2 in zip(I1,I2):
        if (np.random.rand())<p):
            I.append([i1,i2])
    return np.array(I)

def Loss(U,W,Y,I):
    """
    Loss function (2), where
    U, W are n x r factor matrices,
    Y is a n x n data matrix,
    and I is a m x 2 matrix of index pairs
    """
    L = 0
    for i1,i2 in I:
        L = L + (U[i1] @ W[i2] - Y[i1,i2])**2
    return L / I.shape[0]

def initial(n,r):
    """
    Initial guess of the matrix factors
```

```

    """
    U0 = np.vstack((np.identity(r), np.zeros((n-r,r))))
    W0 = np.vstack((np.identity(r), np.zeros((n-r,r))))
    return U0, W0

def test_pointwise_gradient_1():
    """
    Unit tests for the pointwise gradient
    """
    Y = syntheticY(4)
    U,W = initial(4,2)
    vu = np.array([0,-1.125])
    vw = np.array([-1.125,0])
    assert np.array_equal(pointwise_gradient(U,W,Y,0,1), [vu, vw])

def test_pointwise_gradient_2():
    Y = syntheticY(4)
    U,W = initial(4,2)
    vu = np.array([0,0])
    vw = np.array([-0.5,0])
    assert np.array_equal(pointwise_gradient(U,W,Y,0,2), [vu, vw])

def test_pointwise_gradient_3():
    Y = syntheticY(4)
    U,W = initial(4,2)
    vu = np.array([-0.5,0])
    vw = np.array([0,0])
    assert np.array_equal(pointwise_gradient(U,W,Y,2,0), [vu, vw])

def test_full_gradient_1():
    """
    Unit tests for the full gradient
    """
    Y = syntheticY(4)
    U,W = initial(4,2)
    Gu = np.array([[0,-1.125], [0, 0], [0,0], [0,0]])
    Gw = np.array([[0, 0], [-1.125,0], [0,0], [0,0]])
    assert np.array_equal(full_gradient(U,W,Y, np.array([[0,1]])), [Gu, Gw])

def test_full_gradient_2():
    Y = syntheticY(4)
    U,W = initial(4,2)
    Gu = np.array([[0, 0], [-9/16,0], [0, 0], [0,0]])
    Gw = np.array([[0,-9/16], [0, 0], [-1/4,0], [0,0]])
    assert np.array_equal(full_gradient(U,W,Y, np.array([[1,0], [0,2]])),  
↪ [Gu, Gw])

def test_full_gradient_3():
    Y = syntheticY(4)

```

```

U,W = initial(4,2)
Gu = np.array([[0,-9/32], [-9/32,0], [-1/8,0], [0,0]])
Gw = np.array([[0,-9/32], [-9/32,0], [0, 0], [0,0]])
assert np.array_equal(full_gradient(U,W,Y, np.array([[1,0], [0,1],
↪ [2,0], [0,0]])), [Gu, Gw])

```

## 1 Question 1

✓

```

[2]: def pointwise_gradient(U,W,Y,i1,i2):
    # compute formulas (5) and (6) in the assignment
    vu = 2*(U[i1,:] @ W[i2,:].T - Y[i1,i2])* W[i2,:]
    vw = 2*(U[i1,:] @ W[i2,:].T - Y[i1,i2])* U[i1,:]
    return vu,vw

```

## 2 Question 2

✓

```

[3]: def full_gradient(U,W,Y,I):
    # initialise Gu and Gw to same dimensions as U and W
    Gu = np.zeros((U.shape[0],U.shape[1]))
    Gw = np.zeros((U.shape[0],U.shape[1]))
    # iterate over samples
    for i in range(I.shape[0]):
        i1 = I[i,0]
        i2 = I[i,1]
        # compute formulas (5) and (6) in the assignment
        [vu,vw] = pointwise_gradient(U,W,Y,i1,i2)
        # increment ith row of Gu and Gw
        Gu[i1,:] = Gu[i1,:]+vu
        Gw[i2,:] = Gw[i2,:]+vw
    Gu=Gu/I.shape[0]
    Gw=Gw/I.shape[0]
    return Gu, Gw

```

## 3 Question 3

```

[4]: from matplotlib import pyplot as plt
def gd(U0, W0, Y, I, Itest, t=1,eps=1e-6, K=100):
    #initialise starting values and arrays for iterations
    U = U0
    W = W0
    LossVals = np.zeros(0)
    kdagg = Loss(U,W,Y,Itest)
    for k in range(K):
        # save loss for log plot
        LossVals = np.append(LossVals, Loss(U,W,Y,Itest))
        # is loss sufficiently accurate?
        if (LossVals[k] < eps):
            print("converged in "+str(k)+" iterations")

```

```

        # if so exit iteration loop
        break
    if (LossVals[k] > 100):
        print('loss value at iteration '+str(k)+' is too big, stop_
↳algorithm t='+str(t))
        break
    # if not then calculate a sufficient update for U and W
    Gu,Gw = full_gradient(U,W,Y,I)
    U = U - t*Gu
    W = W - t*Gw
    print('the loss value at '+str(k)+' was '+str(LossVals[k])+' for_
↳t='+str(t))
    karray = np.arange(0,k+1)
    # plot k iterations on x-axis and loss values for each k on y-axis
    plt.plot(karray, LossVals, label = 't='+str(t)+'')
    # set the y-axis to log scale
    plt.yscale('log')
    # return U and W which minimise the loss function to sufficient accuracy
    return U,W

```

## 4 Question 4

Ran 3 times and then computed the average for reproduceability (Note: using different testing and training sets for each run)

```

[5]: # function to create I, Itest
def createIs(Iext):
    # shuffle an array of indicies of length Iext
    ind = np.arange(Iext.shape[0])
    np.random.shuffle(ind)
    # shuffle data pairs according to suffled indexes
    Irand = Iext[ind]
    # select first mu pairs to be the testing data set
    Itest = Irand[:mu,:]
    # select the remaining pairs to be the I data set
    I = Irand[Itest.shape[0]:Irand.shape[0]]
    return I, Itest

```

```

[6]: # initialising all variables required in description of question

np.random.seed(0)
n=32
p=1
r = 2
mu = 20
Y = syntheticY(n)
Iext = bernoulli(n,p)
U0,W0=initial(n,r)
I, Itest = createIs(Iext)
K = 5000


```

```
eps = 1e-8
t = [1,2,4,8,16,32]
```


```
[7]: # 1st run

# plotting log of loss against iterations k
frobnormsave = []
for i in t:
    U,W = gd(U0, W0, Y, I, Itest, i,eps, K)
    # calcuting the frobenius norm and saving it for each t
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave = np.append(frobnormsave,frobnorm)
# titles for plots
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Gradient Descent (GD) method for Matrix Completion (run 1)')
plt.legend()
plt.show()
# plotting t = 64 seperately due to different resolution
U,W = gd(U0, W0, Y, I, Itest, 64,eps, K)
frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
frobnormsave = np.append(frobnormsave,frobnorm)
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Gradient Descent (GD) method for Matrix Completion (run 1)')
plt.legend()
plt.show()
```

converged in 1120 iterations  
the loss value at 1120 was 9.980461676684319e-09 for t=1  
converged in 560 iterations  
the loss value at 560 was 9.853979004444943e-09 for t=2  
converged in 280 iterations  
the loss value at 280 was 9.60203115704594e-09 for t=4  
converged in 140 iterations  
the loss value at 140 was 9.102122493833053e-09 for t=8  
converged in 70 iterations  
the loss value at 70 was 8.116977336146316e-09 for t=16  
converged in 35 iterations  
the loss value at 35 was 6.190982801618545e-09 for t=32



loss value at iteration 80 is too big, stop algorithm t=64  
the loss value at 80 was 179.00657256646565 for t=64




```
[8]: # second run
I, Itest = createIs(Iext)
frobnormsave1 = []
```

```


for i in t:
    U,W = gd(U0, W0, Y, I, Itest, i,eps, K)
    # calcuting the frobenius norm and saving it for each t
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave1 = np.append(frobnormsave,frobnorm)
# titles for plots
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Gradient Descent (GD) method for Matrix Completion (run 2)')
plt.legend()
plt.show()
# plotting t = 64 seperately due to different resolution
U,W = gd(U0, W0, Y, I, Itest, 64,eps, K)
frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
frobnormsave1 = np.append(frobnormsave,frobnorm)
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Gradient Descent (GD) method for Matrix Completion (run 2)')
plt.legend()
plt.show()

```

converged in 989 iterations  
 the loss value at 989 was 9.895551369911352e-09 for t=1  
 converged in 495 iterations  
 the loss value at 495 was 9.63206813686197e-09 for t=2  
 converged in 247 iterations  
 the loss value at 247 was 9.926279041326018e-09 for t=4  
 converged in 124 iterations  
 the loss value at 124 was 8.888184955487948e-09 for t=8  
 converged in 62 iterations  
 the loss value at 62 was 8.41468041533058e-09 for t=16  
 converged in 31 iterations  
 the loss value at 31 was 7.421879312266183e-09 for t=32



loss value at iteration 90 is too big, stop algorithm t=64  
the loss value at 90 was 5951762.95178904 for t=64



```
[9]: # third run
I, Itest = createIs(Iext)
frobnormsave2 = []
```



```


for i in t:
    U,W = gd(U0, W0, Y, I, Itest, i,eps, K)
    # calcuting the frobenius norm and saving it for each t
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave2 = np.append(frobnormsave,frobnorm)
# titles for plots
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Gradient Descent (GD) method for Matrix Completion (run 3)')
plt.legend()
plt.show()
# plotting t = 64 seperately due to different resolution
U,W = gd(U0, W0, Y, I, Itest, 64,eps, K)
frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
frobnormsave2 = np.append(frobnormsave,frobnorm)
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Gradient Descent (GD) method for Matrix Completion (run 3)')
plt.legend()
plt.show()

```


```

converged in 1274 iterations
the loss value at 1274 was 9.855748448498271e-09 for t=1
converged in 637 iterations
the loss value at 637 was 9.882654551517296e-09 for t=2
converged in 319 iterations
the loss value at 319 was 9.602505517800192e-09 for t=4
converged in 160 iterations
the loss value at 160 was 9.051379183888803e-09 for t=8
converged in 80 iterations
the loss value at 80 was 9.232691054394484e-09 for t=16
converged in 40 iterations
the loss value at 40 was 9.557535793971314e-09 for t=32

```



loss value at iteration 119 is too big, stop algorithm t=64  
the loss value at 119 was 692005.8325373644 for t=64



Average Frobenius norm across all runs for written pdf.

```
[11]: print('Average Frobenius norm for increasing t^i from 1-64')
      for i in range(7):
          print((1/3)*(frobnormsave[i] + frobnormsave1[i] + frobnormsave2[i]))
```

```
0.0001896422060749921
0.00018777262845073606
0.00018398432435243339
0.0001762156444120005
0.0001599516027254508
0.0001249989603203094
48.63534797649846
```

## 5 Question 5

```
[16]: def sag(U0, W0, Y, I, Itest, t=1, eps=1e-6, K=100):
        #initialise matrices
        U = U0
        W = W0
        m = I.shape[0]
        Vu_tilde = np.zeros((m,U.shape[1]))
        Vw_tilde = np.zeros((m,U.shape[1]))
        Gu_tilde = np.zeros((U.shape[0],U.shape[1]))
        Gw_tilde = np.zeros((U.shape[0],U.shape[1]))
        LossValsS = np.zeros(0, dtype=float)
        for k in range(K):
            LossValsS = np.append(LossValsS, Loss(U,W,Y,Itest))
            if (LossValsS[k] < eps):
                print("converged in "+str(k)+" iterations")
                # if so exit iteration loop
                break
            if (LossValsS[k] > 100):
                print('loss value at iteration '+str(k)+' is too big, stop_
↪algorithm for t='+str(t))
                break
            # sample i uniformly at random
            i = np.arange(m)
            np.random.shuffle(i)
            i = i[0]
            # calculate the corresponding ith pair of I using random i
            i1 = I[i,0]
            i2 = I[i,1]
            # subtract old gradients
            Gu_tilde[i1,:] = Gu_tilde[i1,:] - (1/m)*Vu_tilde[i,:]
            Gw_tilde[i2,:] = Gw_tilde[i2,:] - (1/m)*Vw_tilde[i,:]
            # compute new gradients
            [Vu_tilde[i,:], Vw_tilde[i,:]] = pointwise_gradient(U,W,Y,i1,i2)
            # Add new gradients
            Gu_tilde[i1,:] = Gu_tilde[i1,:] + (1/m)*Vu_tilde[i,:]
            Gw_tilde[i2,:] = Gw_tilde[i2,:] + (1/m)*Vw_tilde[i,:]
            # Update iterates
```

✓

```

        U = U-t*Gu_tilde
        W = W-t*Gw_tilde
        # create an array for the k iterations
        karray = np.arange(0,k+1)
        # plot k iterations on x-axis and loss values for each k on y-axis
        plt.plot(karray, LossValsS, label = 't='+str(t)+'')
        # set the y-axis to log scale
        plt.yscale('log')
        print('the loss value at '+str(k)+' was '+str(LossValsS[k])+' for_
↪t='+str(t))
        # return U and W which minimise the loss function to sufficient accuracy
        return U,W

```

## 6 Question 6

Ran 3 times for reproducability in same procedure as question 4, with different maximum iterations

✓ [33]: #the number of iterations is changed from question 4

```

K = 50000
eps = 1e-8
# plotting 1, 1/2 seperately due to different resolution

# RUN 1
# create a different permutation for I and Itest
I, Itest = createIs(Iext)
frobnormsave = []
t = [1,1/2]
for i in t:
    U,W = sag(U0, W0, Y, I, Itest, i,eps, K)
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave = np.append(frobnormsave,frobnorm)
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Stochastic Average Gradient (SAG) method for Matrix Completion_
↪(run 1)')
plt.legend()
plt.show()
# plotting t values that have similar resolution that dont blow up!
t = [1/4,1/8,1/16,1/32,1/64]


for i in t:
    U,W = sag(U0, W0, Y, I, Itest, i,eps, K)
    # calcuting the frobenius norm and saving it for each t
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave = np.append(frobnormsave,frobnorm)
# titles for plots
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')

```

```
plt.title('Stochastic Average Gradient (SAG) method for Matrix Completion_↵
↵(run 1)')
plt.legend()
plt.show()
```


loss value at iteration 3012 is too big, stop algorithm for  $t=1$   
the loss value at 3012 was 104.47399315541789 for  $t=1$   
loss value at iteration 7427 is too big, stop algorithm for  $t=0.5$   
the loss value at 7427 was 102.24347451991909 for  $t=0.5$

Stochastic Average Gradient (SAG) method for Matrix Completion (run 1)



the loss value at 49999 was 1.3622219276664165e-07 for  $t=0.25$   
converged in 34397 iterations  
the loss value at 34397 was 9.997063446770463e-09 for  $t=0.125$   
converged in 31923 iterations  
the loss value at 31923 was 9.987824776294165e-09 for  $t=0.0625$   
converged in 33684 iterations  
the loss value at 33684 was 9.994760637315859e-09 for  $t=0.03125$   
the loss value at 49999 was 2.427565972114296e-05 for  $t=0.015625$


Stochastic Average Gradient (SAG) method for Matrix Completion (run 1)



```
[34]: # RUN 2
# create a different permutation for I and Itest
I, Itest = createIs(Iext)
frobnormsave1 = []
t = [1,1/2]
for i in t:
    U,W = sag(U0, W0, Y, I, Itest, i,eps, K)
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave1 = np.append(frobnormsave,frobnorm)
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Stochastic Average Gradient (SAG) method for Matrix Completion_
↪(run 2)')
plt.legend()
plt.show()
# plotting t values that have similar resolution that dont blow up!
t = [1/4,1/8,1/16,1/32,1/64]
for i in t:
    U,W = sag(U0, W0, Y, I, Itest, i,eps, K)
    # calcuting the frobenius norm and saving it for each t
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave1 = np.append(frobnormsave,frobnorm)
# titles for plots
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Stochastic Average Gradient (SAG) method for Matrix Completion_
↪(run 2)')
plt.legend()
plt.show()
```


loss value at iteration 2918 is too big, stop algorithm for  $t=1$   
the loss value at 2918 was 104.1794905692511 for  $t=1$   
loss value at iteration 10027 is too big, stop algorithm for  $t=0.5$   
the loss value at 10027 was 124.38840082334471 for  $t=0.5$

Stochastic Average Gradient (SAG) method for Matrix Completion (run 2)



the loss value at 49999 was  $7.064275549208246 \times 10^{-7}$  for  $t=0.25$   
converged in 34510 iterations  
the loss value at 34510 was  $9.99179256299337 \times 10^{-9}$  for  $t=0.125$   
converged in 33303 iterations  
the loss value at 33303 was  $9.99564911667753 \times 10^{-9}$  for  $t=0.0625$   
converged in 40340 iterations  
the loss value at 40340 was  $9.99986032126327 \times 10^{-9}$  for  $t=0.03125$   
the loss value at 49999 was  $0.0002754742488160346$  for  $t=0.015625$

Stochastic Average Gradient (SAG) method for Matrix Completion (run 2)




```
[36]: # RUN 3
# create a different permutation for I and Itest
I, Itest = createIs(Itest)
frobnormsave2 = []
t = [1,1/2]
for i in t:
    U,W = sag(U0, W0, Y, I, Itest, i,eps, K)
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave2 = np.append(frobnormsave,frobnorm)
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Stochastic Average Gradient (SAG) method for Matrix Completion_
↪(run 3)')
plt.legend()
plt.show()
# plotting t values that have similar resolution that dont blow up!
t = [1/4,1/8,1/16,1/32,1/64]
for i in t:
    U,W = sag(U0, W0, Y, I, Itest, i,eps, K)
    # calcuting the frobenius norm and saving it for each t
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave2 = np.append(frobnormsave,frobnorm)
# titles for plots
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Stochastic Average Gradient (SAG) method for Matrix Completion_
↪(run 3)')
plt.legend()
plt.show()
```




loss value at iteration 3513 is too big, stop algorithm for  $t=1$   
the loss value at 3513 was 104.94929343487327 for  $t=1$   
loss value at iteration 9815 is too big, stop algorithm for  $t=0.5$   
the loss value at 9815 was 110.90323747086575 for  $t=0.5$

Stochastic Average Gradient (SAG) method for Matrix Completion (run 3)



the loss value at 49999 was 5.209258271153829e-08 for  $t=0.25$   
converged in 34545 iterations  
the loss value at 34545 was 9.99776905129805e-09 for  $t=0.125$   
converged in 37125 iterations  
the loss value at 37125 was 9.998713746375011e-09 for  $t=0.0625$   
converged in 43394 iterations  
the loss value at 43394 was 9.990650752428566e-09 for  $t=0.03125$   
the loss value at 49999 was 0.032885275624277566 for  $t=0.015625$

Stochastic Average Gradient (SAG) method for Matrix Completion (run 3)



✓

```
[38]: print('Average Frobenius norm for increasing i with 1/t^i from 1 to 1/64')
      for i in range(7):
          print((1/3)*(frobnormsave[i] + frobnormsave1[i] + frobnormsave2[i]))
```

```
Average Frobenius norm for increasing i with 1/t^i from 1 to 1/64
30.030491415878288
4251.20587331164
0.0008124153580489851
0.00033638404888243603
0.00024946534498733946
0.00021932800859761044
0.00951639951635317
```

## 7 Question 7

This will be done in steps: \* Loading in the data \* Initialising variables specified in question  
 \* Running GD and SGD x2 and finding  $L_{Low}$  for respective runs \* Finding the  $t$  value which gives the fastest convergence (x2)

```
[17]: # initialising all variables required in description of question
```

```
#loading the data and storing it as a matrix called Y
np.random.seed(1)

data=np.load('CW.npz')
Y=data['Y']
n=Y.shape[0]
mu = 20
p=1/3
```

✓

```

r=4
Iext = bernoulli(n,p)
U0,W0=initial(n,r)
I, Itest = createIs(Iext)
K = 5000
eps = 1.0e-8
t = [1,2,4,8,16,32,64]

```

Running GD two times to check for reducability in results (Note: using different testing and training sets)

```

[51]: ## NOTE: THIS TAKES 25 MINS TO RUN!

# Run 1 of GD
frobnormsave = []
t = [1,2,4,8,16,32]
for i in t:
    U,W = gd(U0, W0, Y, I, Itest, i,eps, K)
    # calcuting the frobenius norm and saving it for each t
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave = np.append(frobnormsave,frobnorm)
# titles for plots
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Gradient Descent (GD) method for Matrix Completion (run 1)')
plt.legend()
plt.show()
# plotting 64 seperately due to different resolution
U,W = gd(U0, W0, Y, I, Itest, 64,eps, K)
frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
frobnormsave = np.append(frobnormsave,frobnorm)
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Gradient Descent (GD) method for Matrix Completion (run 1)')
plt.legend()
plt.show()

# Run 2 of GD
# change I and Itest data sets
I1, Itest1 = createIs(Iext)
frobnormsave1 = []
t = [1,2,4,8,16,32]
for i in t:
    U,W = gd(U0, W0, Y, I1, Itest1, i,eps, K)
    # calcuting the frobenius norm and saving it for each t
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave1 = np.append(frobnormsave,frobnorm)
# titles for plots
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Gradient Descent (GD) method for Matrix Completion (run 2)')

```

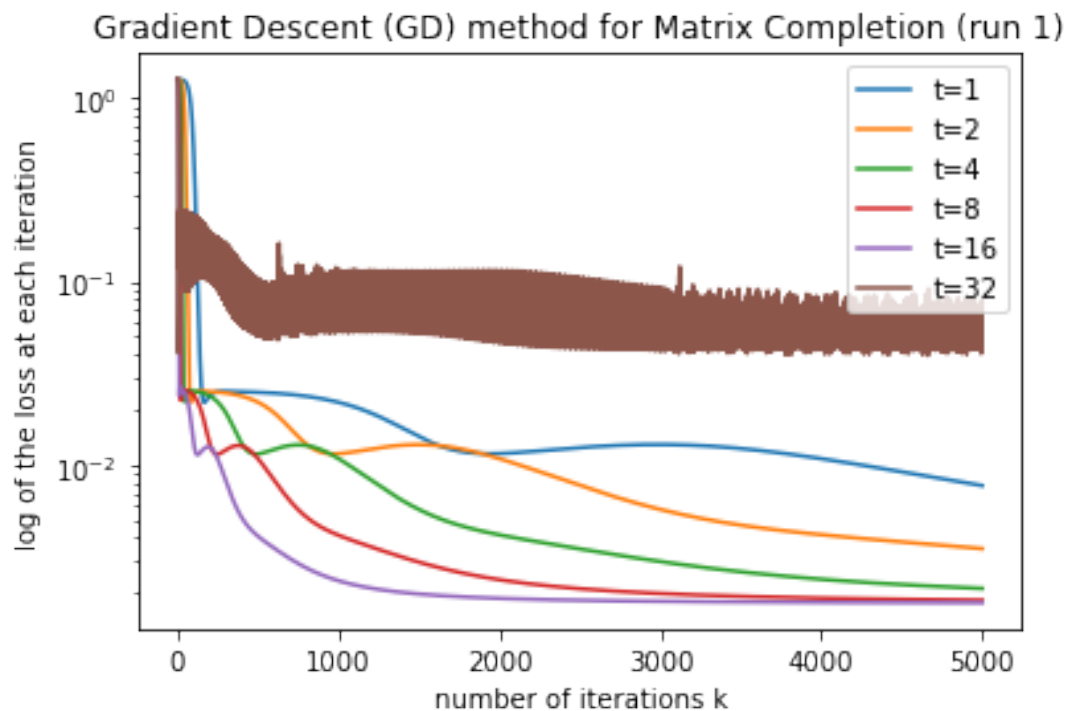
```

plt.legend()
plt.show()
# plotting 64 separately due to different resolution
U,W = gd(U0, W0, Y, I1, Itest1, 64,eps, K)
frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
frobnormsave1 = np.append(frobnormsave,frobnorm)
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Gradient Descent (GD) method for Matrix Completion (run 2)')
plt.legend()
plt.show()

# calculating the average Frobenius norm values for increasing t
print('Average Frobenius norm for increasing t^i from 1-64')
for i in range(7):
    print((1/2)*(frobnormsave[i] + frobnormsave1[i]))


```

the loss value at 4999 was 0.007716006328542054 for t=1  
the loss value at 4999 was 0.0034954883750474295 for t=2  
the loss value at 4999 was 0.002121139037511279 for t=4  
the loss value at 4999 was 0.0018294204513026788 for t=8  
the loss value at 4999 was 0.0017730281702907648 for t=16  
the loss value at 4999 was 0.07910263113667576 for t=32




loss value at iteration 6 is too big, stop algorithm t=64  
the loss value at 6 was 146.33155631767244 for t=64

Gradient Descent (GD) method for Matrix Completion (run 1)

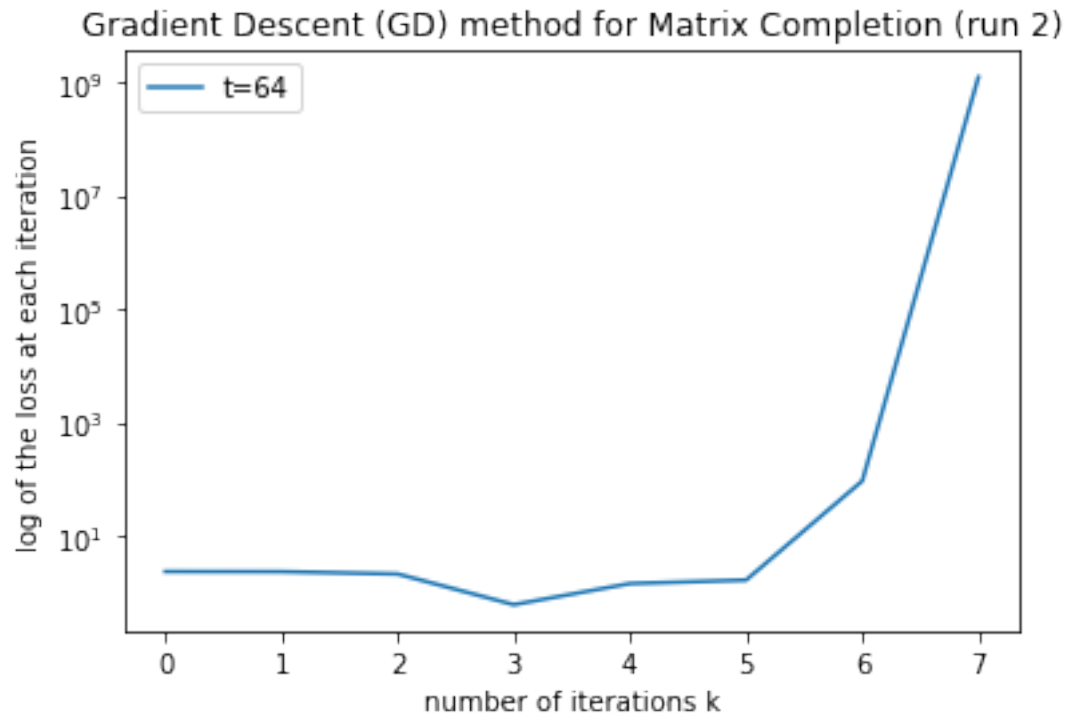


the loss value at 4999 was 0.006460579511994191 for t=1  
the loss value at 4999 was 0.004180110885865918 for t=2  
the loss value at 4999 was 0.005159881789417569 for t=4  
the loss value at 4999 was 0.004224643562368216 for t=8  
the loss value at 4999 was 0.00925924386519562 for t=16  
the loss value at 4999 was 0.23430910573206884 for t=32

Gradient Descent (GD) method for Matrix Completion (run 2)



loss value at iteration 7 is too big, stop algorithm t=64  
the loss value at 7 was 1210047796.0507398 for t=64



Average Frobenius norm for increasing  $t^i$  from 1-64  
0.05472871773064081  
0.0436990171453183  
0.044913686277408896  
0.04557739892247233  
0.04575025533508791  
0.9583462110744541  
38.51327187418235

Running SAG two times to check for reducability in results (Note: using different testing and training sets)

```
[53]: # change number of iterations for this test
K = 50000
eps = 1e-8
# plotting 1, 1/2, 1/4, 1/8 seperately due to different resolution

# RUN 1 of SAG
frobnormsave = []
t = [1, 1/2, 1/4, 1/8]
for i in t:
    U, W = sag(U0, W0, Y, I, Itest, i, eps, K)
    frobnorm = np.linalg.norm(U @ W.T - Y) / np.linalg.norm(Y)
```

```

        frobnormsave = np.append(frobnormsave, frobnorm)
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Stochastic Average Gradient (SAG) method for Matrix Completion,
↳(run 1)')
plt.legend()
plt.show()
# plotting t values that have similar resolution that dont blow up!
t = [1/16, 1/32, 1/64]
for i in t:
    U, W = sag(U0, W0, Y, I, Itest, i, eps, K)
    # calculating the frobenius norm and saving it for each t
    frobnorm = np.linalg.norm(U @ W.T - Y) / np.linalg.norm(Y)
    frobnormsave = np.append(frobnormsave, frobnorm)
# titles for plots
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Stochastic Average Gradient (SAG) method for Matrix Completion,
↳(run 1)')
plt.legend()
plt.show()

# RUN 2 of SAG

frobnormsave1 = []
t = [1, 1/2, 1/4, 1/8]
for i in t:
    U, W = sag(U0, W0, Y, I1, Itest1, i, eps, K)
    frobnorm = np.linalg.norm(U @ W.T - Y) / np.linalg.norm(Y)
    frobnormsave1 = np.append(frobnormsave, frobnorm)
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Stochastic Average Gradient (SAG) method for Matrix Completion,
↳(run 2)')
plt.legend()
plt.show()
# plotting t values that have similar resolution that dont blow up!
t = [1/16, 1/32, 1/64]
for i in t:
    U, W = sag(U0, W0, Y, I1, Itest1, i, eps, K)
    # calculating the frobenius norm and saving it for each t
    frobnorm = np.linalg.norm(U @ W.T - Y) / np.linalg.norm(Y)
    frobnormsave1 = np.append(frobnormsave, frobnorm)
# titles for plots
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Stochastic Average Gradient (SAG) method for Matrix Completion,
↳(run 2)')
plt.legend()
plt.show()


```

```
# compute the average frobenius norm
```

```
print('Average Frobenius norm for increasing  $t^i$  from  $1-1/64$ ')
for i in range(7):
    print((1/2)*(frobnormsave[i] + frobnormsave1[i]))
```

loss value at iteration 1815 is too big, stop algorithm for  $t=1$   
the loss value at 1815 was 103.6848892458178 for  $t=1$   
loss value at iteration 3060 is too big, stop algorithm for  $t=0.5$   
the loss value at 3060 was 108.99659077211541 for  $t=0.5$   
loss value at iteration 6324 is too big, stop algorithm for  $t=0.25$   
the loss value at 6324 was 101.11156288481202 for  $t=0.25$   
loss value at iteration 21087 is too big, stop algorithm for  $t=0.125$   
the loss value at 21087 was 104.28025909932894 for  $t=0.125$


Stochastic Average Gradient (SAG) method for Matrix Completion (run 1)



the loss value at 49999 was 0.010504701650727657 for  $t=0.0625$   
the loss value at 49999 was 0.02487297843880441 for  $t=0.03125$   
the loss value at 49999 was 0.02400160767220706 for  $t=0.015625$




Stochastic Average Gradient (SAG) method for Matrix Completion (run 1)




loss value at iteration 1850 is too big, stop algorithm for  $t=1$   
the loss value at 1850 was 151.19790475715965 for  $t=1$   
loss value at iteration 2990 is too big, stop algorithm for  $t=0.5$   
the loss value at 2990 was 100.88114550279529 for  $t=0.5$   
loss value at iteration 6608 is too big, stop algorithm for  $t=0.25$   
the loss value at 6608 was 103.4857859985356 for  $t=0.25$   
loss value at iteration 20409 is too big, stop algorithm for  $t=0.125$   
the loss value at 20409 was 101.58999986196989 for  $t=0.125$

Stochastic Average Gradient (SAG) method for Matrix Completion (run 2)



the loss value at 49999 was 0.012054233821322136 for  $t=0.0625$   
the loss value at 49999 was 0.015759654120829227 for  $t=0.03125$   
the loss value at 49999 was 0.022970488966494073 for  $t=0.015625$

Stochastic Average Gradient (SAG) method for Matrix Completion (run 2)



Average Frobenius norm for increasing  $t^i$  from  $1-1/64$

13.313392580177505  
22.991319575510236  
28.57295953683959  
390057.6869718698  
0.07702741990466157  
0.1005363547446034  
0.11762976857525714

For SG \* Run 1:  $L_{low} = 0.0017730281702907648$  \* Run 2:  $L_{low} = 0.004180110885865918$  ✓

For SAG \* Run 1:  $L_{low} = 0.010504701650727657$  \* Run 2:  $L_{low} = 0.012054233821322136$  too large

Proceed to re-run tests with  $\epsilon = 2 * L_{low}$  and changed maximum iterations where nessaccery

(NOTE: Originally set  $K = 10000$ , but tests took 1hr to run... Setting  $K = 3000$  is sufficient to obtain the quickest converging  $t$ )

[94]: *## NOTE: THIS TAKES A LONG TIME TO RUN!*  
K=3000  
*# Run 1 of SG1 for eps = 2\*0.0017730281702907648*  
eps = 2\*0.0017730281702907648  
frobnormsave = []  
t = [1,2,4,8,16,32]  
for i in t:  
 U,W = gd(U0, W0, Y, I, Itest, i,eps, K)  
*# calcuting the frobenius norm and saving it for each t*  
 frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)

```

        frobnormsave = np.append(frobnormsave,frobnorm)
# titles for plots
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Gradient Descent (GD) method for Matrix Completion (run 1)')
plt.legend()
plt.show()
# plotting 64 seperately due to different resolution
U,W = gd(U0, W0, Y, I, Itest, 64,eps, K)
frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
frobnormsave = np.append(frobnormsave,frobnorm)
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Gradient Descent (GD) method for Matrix Completion (run 1)')
plt.legend()
plt.show()

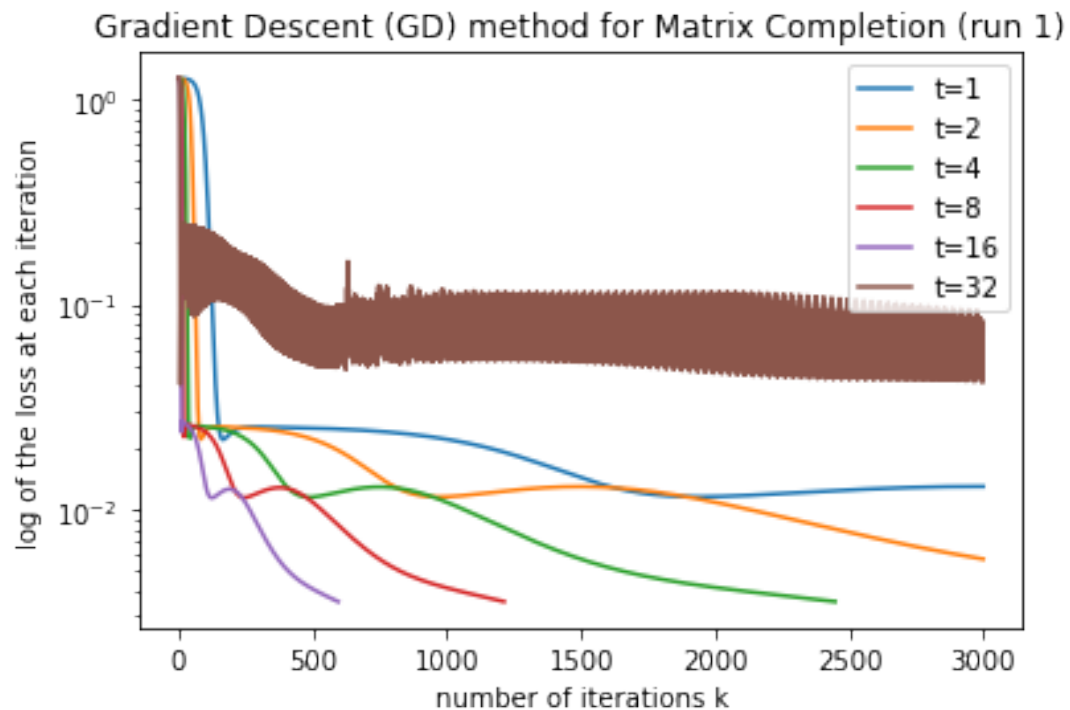
# Run 2 of GD for eps = 2*0.004180110885865918
eps = 2*0.004180110885865918
# change I and Itest data sets
frobnormsave1 = []
t = [1,2,4,8,16,32]
for i in t:
    U,W = gd(U0, W0, Y, I1, Itest1, i,eps, K)
    # calcuting the frobenius norm and saving it for each t
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave1 = np.append(frobnormsave,frobnorm)
# titles for plots
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Gradient Descent (GD) method for Matrix Completion (run 2)')
plt.legend()
plt.show()
# plotting 64 seperately due to different resolution
U,W = gd(U0, W0, Y, I1, Itest1, 64,eps, K)
frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
frobnormsave1 = np.append(frobnormsave,frobnorm)
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Gradient Descent (GD) method for Matrix Completion (run 2)')
plt.legend()
plt.show()

# calculating the average Frobenius norm values for increasing t
print('Average Frobenius norm for increasing t~i from 1-64')
for i in range(7):
    print((1/2)*(frobnormsave[i] + frobnormsave1[i]))


```

the loss value at 2999 was 0.012925027582220552 for t=1  
the loss value at 2999 was 0.005718114045265154 for t=2  
converged in 2447 iterations

the loss value at 2447 was 0.0035457974200612517 for t=4  
 converged in 1213 iterations  
 the loss value at 1213 was 0.003545596602725843 for t=8  
 converged in 595 iterations  
 the loss value at 595 was 0.0035413436630722197 for t=16  
 the loss value at 2999 was 0.08187399806301829 for t=32




loss value at iteration 6 is too big, stop algorithm t=64  
 the loss value at 6 was 146.33155631767244 for t=64




the loss value at 2999 was 0.008956469311753552 for  $t=1$   
converged in 1671 iterations  
the loss value at 1671 was 0.008360026188405274 for  $t=2$   
converged in 853 iterations  
the loss value at 853 was 0.008358337960403642 for  $t=4$   
converged in 443 iterations  
the loss value at 443 was 0.008355574872528677 for  $t=8$   
converged in 236 iterations  
the loss value at 236 was 0.008356740866570738 for  $t=16$   
the loss value at 2999 was 0.20436304556563586 for  $t=32$

Gradient Descent (GD) method for Matrix Completion (run 2)



loss value at iteration 7 is too big, stop algorithm  $t=64$   
the loss value at 7 was 1210047796.0507398 for  $t=64$

Gradient Descent (GD) method for Matrix Completion (run 2)



Average Frobenius norm for increasing  $t^i$  from 1-64

0.06879717947313183  
0.05014335585153289  
0.043728538461155486  
0.04374127219032466  
0.0437817335923266  
0.9325071930242476  
38.51327187418235

```
[71]: K = 50000
# plotting 1, 1/2, 1/4, 1/8 separately due to different resolution

# RUN 1 of SAG with eps = 2*0.010504701650727657
eps = 2*0.010504701650727657
frobnormsave = []
t = [1, 1/2, 1/4, 1/8]
for i in t:
    U, W = sag(U0, W0, Y, I, Itest, i, eps, K)
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave = np.append(frobnormsave, frobnorm)
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Stochastic Average Gradient (SAG) method for Matrix Completion_
↳(run 1)')
plt.legend()
plt.show()
# plotting t values that have similar resolution that dont blow up!
t = [1/16, 1/32, 1/64]
for i in t:
    U, W = sag(U0, W0, Y, I, Itest, i, eps, K)
    # calculating the frobenius norm and saving it for each t
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave = np.append(frobnormsave, frobnorm)
# titles for plots
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Stochastic Average Gradient (SAG) method for Matrix Completion_
↳(run 1)')
plt.legend()
plt.show()

# RUN 2 of SAG with eps = 2*0.012054233821322136
eps = 2*0.012054233821322136
frobnormsave1 = []
t = [1, 1/2, 1/4, 1/8]
for i in t:
    U, W = sag(U0, W0, Y, I1, Itest1, i, eps, K)
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave1 = np.append(frobnormsave, frobnorm)
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
```

```

plt.title('Stochastic Average Gradient (SAG) method for Matrix Completion,
↳(run 2)')
plt.legend()
plt.show()
# plotting t values that have similar resolution that dont blow up!
t = [1/16,1/32,1/64]
for i in t:
    U,W = sag(U0, W0, Y, I1, Itest1, i,eps, K)
    # calcuting the frobenius norm and saving it for each t
    frobnorm = np.linalg.norm(U @ W.T - Y)/np.linalg.norm(Y)
    frobnormsave1 = np.append(frobnormsave,frobnorm)
# titles for plots
plt.ylabel('log of the loss at each iteration')
plt.xlabel('number of iterations k')
plt.title('Stochastic Average Gradient (SAG) method for Matrix Completion,
↳(run 2)')
plt.legend()
plt.show()

# compute the average frobenius norm


print('Average Frobenius norm for increasing t^i from 1-1/64')
for i in range(7):
    print((1/2)*(frobnormsave[i] + frobnormsave1[i]))

```

loss value at iteration 1952 is too big, stop algorithm for t=1  
 the loss value at 1952 was 100.38842439345498 for t=1  
 loss value at iteration 3114 is too big, stop algorithm for t=0.5  
 the loss value at 3114 was 100.14768436538677 for t=0.5  
 loss value at iteration 6655 is too big, stop algorithm for t=0.25  
 the loss value at 6655 was 101.62968008558552 for t=0.25  
 loss value at iteration 15132 is too big, stop algorithm for t=0.125  
 the loss value at 15132 was 100.88923423819332 for t=0.125



Stochastic Average Gradient (SAG) method for Matrix Completion (run 1)




converged in 41158 iterations

the loss value at 41158 was 0.021006828863447274 for  $t=0.0625$

the loss value at 49999 was 0.03630205622143094 for  $t=0.03125$

the loss value at 49999 was 0.0464452649352613 for  $t=0.015625$

Stochastic Average Gradient (SAG) method for Matrix Completion (run 1)




loss value at iteration 1937 is too big, stop algorithm for  $t=1$

the loss value at 1937 was 635.7406648075366 for  $t=1$

loss value at iteration 3186 is too big, stop algorithm for  $t=0.5$


the loss value at 3186 was 100.10698604917111 for  $t=0.5$   
 loss value at iteration 7380 is too big, stop algorithm for  $t=0.25$   
 the loss value at 7380 was 104.82293794831662 for  $t=0.25$   
 loss value at iteration 21348 is too big, stop algorithm for  $t=0.125$   
 the loss value at 21348 was 101.21249541727884 for  $t=0.125$

Stochastic Average Gradient (SAG) method for Matrix Completion (run 2)



converged in 30472 iterations  
 the loss value at 30472 was 0.02410557921888988 for  $t=0.0625$   
 the loss value at 49999 was 0.031830986484798075 for  $t=0.03125$   
 the loss value at 49999 was 0.0482517152961788 for  $t=0.015625$

Stochastic Average Gradient (SAG) method for Matrix Completion (run 2)



Average Frobenius norm for increasing  $t^i$  from 1-1/64

54.25976876148113

14.90388584386282

109083.36090683022

11.371580242420599

0.08473476648252666


0.11226283988133913

0.11784878672134419

computing images  $Y$  and most accurate  $UW^T$  computed. Choose  $eps = 2 * 0.0017730281702907648$  (since this was the minimum loss found overall described above) and  $t = 16$  (quickest average convergence, smallest number of iterations till convergence...) for GD algorithm (accuracy)

```
[96]: K=10000
print('This is the image of Y (IMAGE 1)')
# display matrix Y as image
plt.imshow(Y)
plt.show()
# compute GS with computed eps and sufficient K iterations max
UG,WG = gd(U0, W0, Y, I, Itest, 16,2*0.0017730281702907648, K)
print('This is the image of UW.T computed via Matrix Completion using GD_
      ↪(IMAGE 2)')
# display UW.T as image
plt.imshow(UG @ WG.T)
plt.yscale('linear')
plt.show()
```

This is the image of Y (IMAGE 1)



```
converged in 595 iterations
the loss value at 595 was 0.0035413436630722197 for t=16
This is the image of UW.T computed via Matrix Completion using GD (IMAGE 2)

/tmp/ipykernel_1030203/629753498.py:10: UserWarning: Attempted to set non-
positive top ylim on a log-scaled axis.
Invalid limit will be ignored.
plt.imshow(UG @ WG.T)
```

